

Block statistics for real-time-targeted difficulty adjustments, with absolutely scheduled exponentially rising targets (RTT-ASERT)

Mark B. Lundeborg

November 6, 2019

Consider a real time difficulty algorithm, where the block target ($= 2^{224}/\text{block difficulty}$) for the block with height N and timestamp t_N is defined as:

$$\text{target}_N = \text{target}_0 \exp([t_N - \tilde{t}_N]/\lambda), \quad (1)$$

where $\tilde{t}_N = t_0 + NT$ and where T is the targeted average block interval (e.g., $T \approx 10$ minutes). The parameter λ is a relaxation time and practically could be for instance $6T$ or $12T$. The base target target_0 and genesis time t_0 are not particularly important in this analysis since they just produce offsets in time.

1 Statistics of the block-finding process

At a given time t , a hashrate $H(t)$ is allocated to mining blocks, and furthermore let's assume that all miners have highly synchronized clocks, and are trying to mine a block N with timestamp $t_N = t$ and they immediately publish any found blocks (whether these actually happen in reality is a separate question!). Then the rate probability of producing block N at timestamp t (supposing it has not already been found yet) is given by:

$$\lambda(t) = H(t) \cdot 2^{-256} \text{target}(t).$$

Following wikipedia's formulas on survival analysis, you just need to integrate $\lambda(t)$ to find the "survival function" $S(t)$, which is the complementary

cumulative distribution function of the time when the block is found. Let's say that at some start time t_s , the miners haven't found the block yet, and you want to know the statistics of the predicted *solve time* $\delta = t_N - t_s$. This is

$$S(\delta) = \exp \left(- \int_{t_s}^{t_s + \delta} \lambda(t) dt \right).$$

Now, let's consider the case of $H(t) = H$ constant, and plug in our exponentially varying target Eq. 1. This gives:

$$S(\delta) = \exp \left(-\lambda H 2^{-256} \text{target}_0 \exp([t_s - \tilde{t}_N]/\lambda) (\exp(\delta/\lambda) - 1) \right) \quad (2)$$

This looks a bit hairy, but let's separate the big constant factor as X :

$$X = \lambda H 2^{-256} \text{target}_0 \exp([t_s - \tilde{t}_N]/\lambda), \quad (3)$$

$$S(\delta) = \exp \left(-X (\exp(\delta/\lambda) - 1) \right). \quad (4)$$

This dimensionless X can be thought of as the ‘‘shape parameter’’. For large $X \gg 1$, this will be the familiar $S(\delta) \approx \exp(-X\delta/\lambda)$ with a maximum initial slope $-X/\lambda$ occurring at $\delta = 0$, i.e., as if the target is independent of time, since the block will be found ‘‘fast’’ relative to time λ . With small $X \ll 1$ in contrast, $S(\delta) \approx \exp[-\exp(\delta/\lambda + \ln X)]$ which has a more sigmoid shape, with a crossing near $\delta \approx \lambda \ln(1/X)$, a maximum slope of around $-0.37/\lambda$, but a rather hard cutoff for 2λ beyond this crossing.

2 Jump-forward simulation

In general, mining simulations can be naively done in a ‘real time’ nature by sampling $\lambda(t)$ at very small time intervals, however these simulations will be slow and thus in a given time will not generate as much useful information. When $H(t)$ is known ahead of time, it is *much* more efficient to do ‘jump-forward’ simulations by simply calculating a random δ with appropriate distribution. Since S is the complementary cumulative distribution function of the block event, you just need to set $S = x$, where x is a random number chosen uniformly between 0 and 1. Then solve for δ .

Taking the constant- H case considered above in Eq. 4, this yields

$$\delta = \lambda \ln(1 - (\ln x)/X) \quad (5)$$

where again the X is a constant defined above. This is slightly more complex than the formula used in static target simulations, but not much more so.

Switch mining simulations typically consider the case where $H(t)$ is piecewise-constant: a low value $H(t) = H_a$ when the target is below some critical value, $\text{target}(t) < \text{target}_{\text{crit}}$, then jumping to a high value $H(t) = H_b$ afterwards. In simulations this is easiest to achieve by calculating δ_{crit} , the time of switch, then generating a δ assuming H_a . If $\delta < \delta_{\text{crit}}$ then the block was found, otherwise jump the simulation forward in time by δ_{crit} (i.e., $t_s \rightarrow t_s + \delta_{\text{crit}}$), then calculate a new X for the new t_s and H_b , and calculate a new random δ measured from the new t_0 . Arbitrary piecewise-constant $H(t)$ can be done in a jump-forward simulation using this trick, though of course for each jump it requires knowing the exact time of the next change in hashrate.

3 Stochastic state properties

In practice, the feedback nature of difficulty targeting means that δ will seek towards a typical value of T , the target block interval. Since slow relaxation with $\lambda > T$ ought to be chosen, this means the chain will primarily be operating in the $X \approx \lambda/T > 1$ regime. Statistical fluctuations and temporary high, sustained hashpower increases can push down X , but the system quickly recovers on a λ timescale.

Since the targeting equation is so simple (it doesn't depend on prior blocks' timestamps nor their targets), at any given time the state of the system can be represented by a single variable: the current target. The target exponentially rises and every time a block is found, the target changes by a factor of $\exp(-T/\lambda)$.

An even simpler picture of dynamics arises when we take logarithm of target as our variable. For convenience let's work in $b = \ln(2^{256}/\text{target})$ space, which is basically the log of difficulty (taking difficulty=1 to mean a block can be solved in 1 hash). So, this means the b ramps *linearly* down over time, at a rate of $1/\lambda$. Then for every block found, b jumps up by an offset $+T/\lambda$. Note that at any given time, the target can only take on specific discrete values (depending on the height N), and those values recur every time T . To get a picture of the 'steady state', it is necessary to average the dynamics over a time T period.

I'm not sure how to solve this equation, not even in the steady state (where it becomes a retarded functional differential equation). However the

behaviour is well suited to numerical simulation, where for each time step the $p(b)$ shifts left by one pixel, and then some fraction is shifted right. The steady state distributions are rather interesting looking, for example below with $\lambda = 12T$, where the “steady” hashrate for $b > 0$ is only $0.2/T$, but then for $b < 0$ some switch miners come in and the hashrate is $5/T$. Note, however, that real time block difficulty algorithms don’t even need a steady hashrate, since there is no hazard of ‘freezing up’!

